

# Generalized Linear Models with Stan

Ben Goodrich

April 13, 2022

# What about the States?

- Suppose we wanted to include an intercept for each state, rather than merely an indicator for whether the state has a Republican governor
- We could include 50 dummy variables in  $\mathbf{X}$  and specify priors on those coefficients, but McElreath prefers the following approach

```
X <- as.matrix(Gabba$Trump - mean(Gabba$Trump))  
group <- as.factor(Gabba$State)  
nlevels(group) # size N but only 51 unique values
```

```
## [1] 51
```

- We can also utilize normal priors if we prefer with means and standard deviations as

```
m      <- c(beta = -0.5, alpha = 50, sigma = 10)  
scale  <- c(beta = 0.25, alpha = 10, sigma = 3)
```

```

data {
    // saved as "groups.stan"
    int<lower = 0> N; // number of observations
    int<lower = 0> K; // number of predictors
    matrix[N, K] X;  // matrix of predictors
    vector[N] y;     // outcomes
    int<lower = 1> J; // number of groups
    int<lower = 1, upper = J> group[N]; // group membership
    int<lower = 0, upper = 1> prior_only; // ignore data?
    vector[K + 2] m; // prior means
    vector<lower = 0>[K + 2] scale; // prior scales
}

parameters {
    vector[K] beta;
    vector[J] alpha;
    real<lower = 0> sigma;
}

model {
    if (!prior_only) target += normal_id_glm_lpdf(y | X, alpha[group], beta, sigma);
    target += normal_lpdf(beta | m[1:K], scale[1:K]); // ^ important
    target += normal_lpdf(alpha | m[K + 1], scale[K + 1]);
    target += normal_lpdf(sigma | m[K + 2], scale[K + 2]); // actually half normal
}

generated quantities {
    vector[N] log_lik;
    {
        vector[N] mu = alpha[group] + X * beta;
    }
}

```

# Calling **stan** for the grouped model

```
states <- stan("groups.stan", data = list(N = nrow(Gabba), K = ncol(X), y = Gabba$Vaccinated,  
                                           X = X, J = nlevels(group), group = as.integer(group)  
                                           prior_only = 0, m = m, scale = scale)) # ^ important
```

*states # only 6 states could fit on the screen but all 51 intercepts were estimated*

```
## Inference for Stan model: groups.
```

```
## 4 chains, each with iter=2000; warmup=1000; thin=1;
```

```
## post-warmup draws per chain=1000, total post-warmup draws=4000.
```

```
##
```

##	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%
## beta[1]	-0.46	0.00	0.01	-0.48	-0.47	-0.46	-0.46	-0.44
## alpha[1]	42.90	0.01	0.89	41.16	42.29	42.89	43.51	44.60
## alpha[2]	55.17	0.02	1.45	52.28	54.21	55.21	56.14	58.03
## alpha[3]	57.23	0.03	2.08	53.22	55.80	57.20	58.68	61.27
## alpha[4]	48.17	0.01	0.86	46.51	47.60	48.17	48.75	49.83
## alpha[5]	51.37	0.01	0.97	49.47	50.73	51.35	52.02	53.27
## alpha[6]	51.56	0.01	0.93	49.76	50.92	51.55	52.21	53.39
## alpha[7]	62.80	0.03	2.57	57.79	61.07	62.80	64.56	67.69
## alpha[8]	52.52	0.05	4.01	44.96	49.74	52.57	55.33	60.33
## alpha[9]	48.63	0.08	5.93	37.15	44.65	48.60	52.63	60.45
## alpha[10]	53.12	0.01	0.94	51.22	52.48	53.13	53.76	54.93
## alpha[11]	42.96	0.01	0.59	41.78	42.55	42.95	43.35	44.11

# Utility Function for Predictions of Future Data

- For Bayesians, the log predictive PDF is the most appropriate utility function
- Choose the model that maximizes the expectation of this over FUTURE data

$$\begin{aligned}\text{ELPD} &= \mathbb{E}_Y \ln f(y_{N+1}, \dots, y_{2N} \mid y_1, \dots, y_N) = \\ &\int \ln f(y_{N+1}, \dots, y_{2N} \mid \mathbf{y}) f(y_{N+1}, \dots, y_{2N} \mid \mathbf{y}) dy_{N+1} \dots dy_{2N} \approx \\ &\sum_{n=1}^N \ln f(y_n \mid \mathbf{y}_{-n}) = \sum_{n=1}^N \ln \int_{\Theta} f(y_n \mid \boldsymbol{\theta}) f(\boldsymbol{\theta} \mid \mathbf{y}_{-n}) d\theta_1 d\theta_2 \dots d\theta_K\end{aligned}$$

- $f(y_n \mid \boldsymbol{\theta})$  is just the  $n$ -th likelihood contribution, but can we somehow obtain  $f(\boldsymbol{\theta} \mid \mathbf{y}_{-n})$  from  $f(\boldsymbol{\theta} \mid \mathbf{y})$ ?
- Yes, assuming  $y_n$  does not have an outsized influence on the posterior

# Optional generated quantities Block

- Can declare more endogenous knowns, assign to them, and use them
- Samples are stored
- Can reference anything except stuff in the `model` block
- Can also do this in R afterward, but primarily used for
  - Interesting functions of posterior that don't involve likelihood
  - Posterior predictive distributions and / or functions thereof
  - The log-likelihood for each observation to pass to `loo`

# PSISLOOCV

```
generated quantities { // part of groups.stan
```

```
  vector[N] log_lik;
  {
    vector[N] mu = alpha[group] + X * beta;
    for (n in 1:N) log_lik[n] = normal_lpdf(y[n] | mu[n], sigma);
  }
}
```

```
loo(states)
```

```
## Warning: Some Pareto k diagnostic values are slightly high. See help('pareto-k-diagnostic')
```

```
##
```

```
## Computed from 4000 by 3125 log-likelihood matrix
```

```
##
```

```
##           Estimate      SE
```

```
## elpd_loo -10788.1  76.7
```

```
## p_loo      51.5   2.8
```

```
## looic      21576.1 153.4
```

```
## -----
```

```
## Monte Carlo SE of elpd_loo is 0.1.
```


```
##
```

```
## Pareto k diagnostic values:
```

# Leverage Diagnostic Plot

```
plot(loo(states), label_points = TRUE) # not too bad, 318 is D.C.
```

```
## Warning: Some Pareto k diagnostic values are slightly high. See help('pareto-k-diagnostic')
```



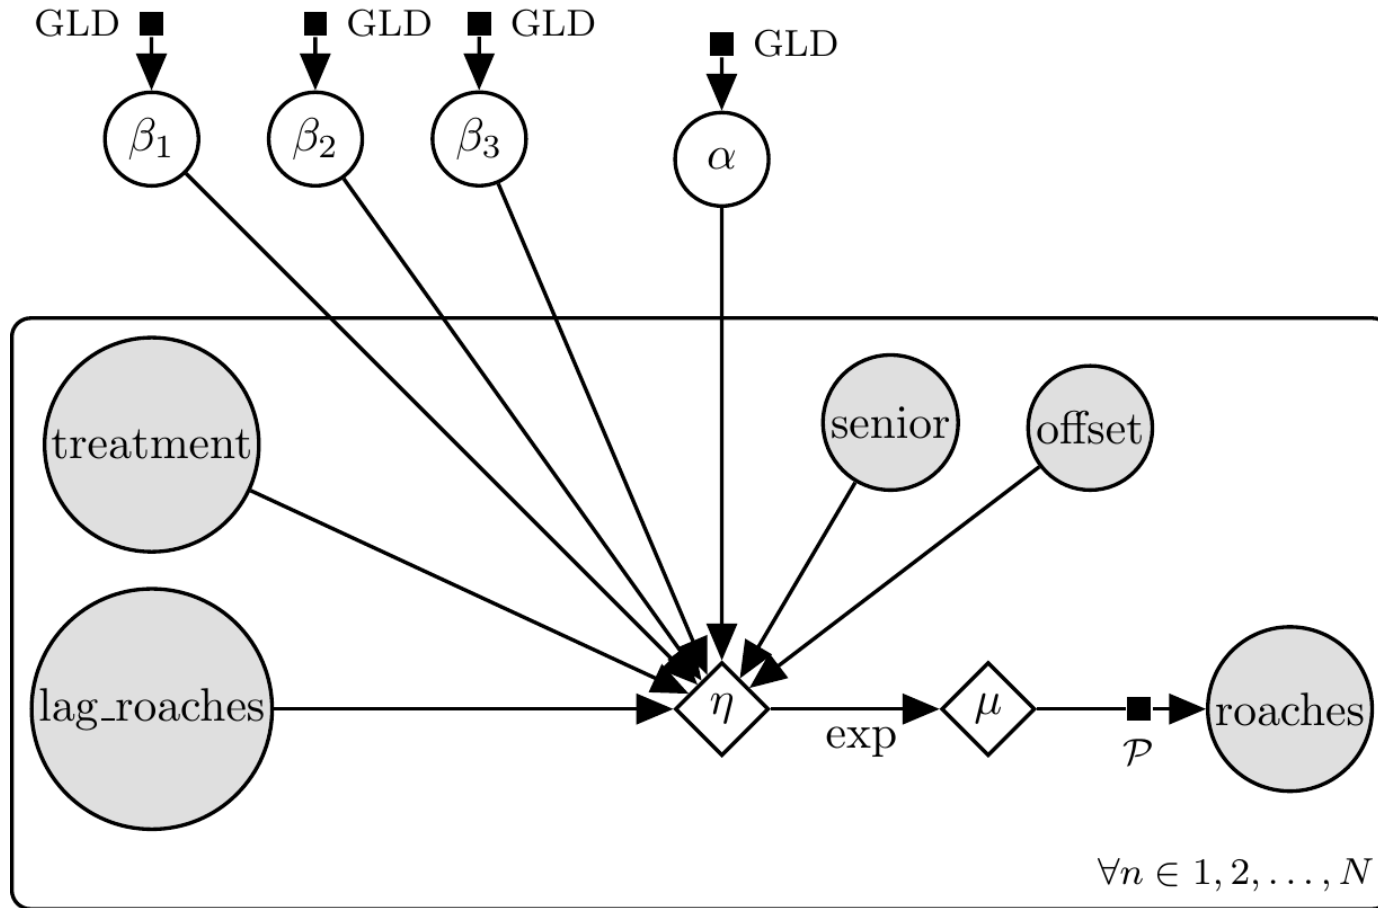


# Roach Data in NYC Experiment

```
ROOT <- "https://raw.githubusercontent.com/avehtari/ROS-Examples/master/"
FILE <- "Roaches/data/roaches.csv"
roaches <- readr::read_csv(paste0(ROOT, FILE), col_types = "_idiid")
library(dplyr)
roaches <- filter(roaches, roach1 > 0)
glimpse(roaches)
```

```
## Rows: 202
## Columns: 5
## $ y      <int> 153, 127, 7, 7, 0, 73, 24, 2, 2, 0, 21, 0, 179, 136, 104, 2, 5, 1, 203...
## $ roach1  <dbl> 308.00, 331.25, 1.67, 3.00, 2.00, 70.00, 64.56, 1.00, 14.00, 138.25, 1...
## $ treatment <int> 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...
## $ senior  <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...
## $ exposure2 <dbl> 0.80000000, 0.60000000, 1.00000000, 1.00000000, 1.1428571, 0.80000000, 1.14...
```

# Prior Predictive Distribution for Roach Study



Roach Model

# Prior Predictive Distribution in Symbols

$$\alpha \sim GLD$$

$$\beta_1 \sim GLD$$

$$\beta_2 \sim GLD$$

$$\beta_3 \sim GLD$$

$$\forall n : \eta_n \equiv \alpha + OFFSET_n + \beta_1 \times \log LAG_n + \beta_2 \times SENIOR_n + \beta_3 \times T_n$$

$$\forall n : \mu_n \equiv e^{\eta_n}$$

$$\forall n : Y_n \sim \mathcal{P}(\mu_n)$$

- In this case, the inverse link function mapping the linear predictor  $\eta_n$  on  $\mathbb{R}$  to the outcome's conditional expectation  $\mu_n$  on  $\mathbb{R}_+$  is the antilog function.
- An “offset” is a predictor whose coefficient is fixed to be 1

# Generalized Lambda Distribution Priors

- What do you believe about  $\beta_1$ , the coefficient on the logarithm of roaches in the previous period?
- What do you believe about  $\beta_2$ , the coefficient on whether the building is a senior living facility?
- What do you believe about  $\beta_3$ , the coefficient on the treatment variable?
- What do you believe about  $\alpha$ , the expected logarithm of roaches for a building with average predictors?

# Calling **stan** for the Poisson Model

```
stan_data <- list(N = nrow(roaches), offset = log(roaches$exposure2),  
                 K = 3, y = roaches$y,  
                 X = with(roaches, cbind(log(roach1) - mean(log(roach1)),  
                                          senior - mean(senior),  
                                          treatment - mean(treatment))),  
                 prior_only = FALSE, m = m, r = r, a = a, s = s)  
  
post_poisson <- stan("poisson.stan", data = stan_data)  
  
## Trying to compile a simple C file
```

# ShinyStan

```
y <- roaches$y  
shinytan::launch_shinytan(post_poisson) # opens in a web browser
```

# Numerical Assessment of Calibration

```
y_rep <- rstan::extract(post_poisson, "y_rep")[[1]]; dim(y_rep)
```

```
## [1] 4000 202
```

```
lower <- apply(y_rep, MARGIN = 2, FUN = quantile, probs = 0.25)
```

```
upper <- apply(y_rep, MARGIN = 2, FUN = quantile, probs = 0.75)
```

```
mean(roaches$y > lower & roaches$y < upper) # bad fit
```

```
## [1] 0.04950495
```

- Overall, the model is fitting the data poorly in this case, although overfitting can be a concern in other situations

# Adding Overdispersion

$$\alpha \sim GLD$$

$$\beta_1 \sim GLD$$

$$\beta_2 \sim GLD$$

$$\beta_3 \sim GLD$$

$$\forall n : \eta_n \equiv \alpha + OFFSET_n + \beta_1 \times \log LAG_n + \beta_2 \times SENIOR_n + \beta_3 \times T_n$$

$$\forall n : \mu_n \equiv e^{\eta_n}$$

$$\phi \sim GLD$$

$$\forall n : \epsilon_n \sim \mathcal{G}(\phi, \phi)$$

$$\forall n : Y_n \sim \text{Poisson}(\epsilon_n \mu_n)$$

\* The conditional distribution of  $Y_n$  given  $\epsilon_n \mu_n$  is Poisson, but the conditional distribution of  $Y_n$  given  $\mu_n$  irrespective of  $\epsilon_n$  is negative binomial with expectation  $\mu_n$  and variance  $\mu_n + \mu_n^2 / \phi$  \* What are your beliefs about  $\phi$ ?



# Calling **stan** for the Negative Binomial Model

```
stan_data$m <- m; stan_data$r <- r; stan_data$a <- a; stan_data$s <- s
```

```
post_nb <- stan("negative_binomial.stan", data = stan_data)
```

```
print(post_nb, pars = c("alpha", "beta", "phi"))
```

```
## Inference for Stan model: negative_binomial.
```

```
## 4 chains, each with iter=2000; warmup=1000; thin=1;
```

```
## post-warmup draws per chain=1000, total post-warmup draws=4000.
```

```
##
```

##		mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
##	alpha	2.82	0.00	0.10	2.61	2.74	2.81	2.88	3.02	4472	1
##	beta[1]	0.71	0.00	0.06	0.58	0.67	0.71	0.75	0.83	4070	1
##	beta[2]	-0.18	0.00	0.24	-0.63	-0.33	-0.18	-0.01	0.28	4086	1
##	beta[3]	-0.49	0.01	0.23	-0.93	-0.65	-0.50	-0.33	-0.01	2118	1
##	phi	0.47	0.00	0.05	0.38	0.44	0.47	0.51	0.58	3870	1

```
##
```

```
## Samples were drawn using NUTS(diag_e) at Wed Apr 13 11:42:31 2022.
```

```
## For each parameter, n_eff is a crude measure of effective sample size,
```

```
## and Rhat is the potential scale reduction factor on split chains (at
```

```
## convergence, Rhat=1).
```

# Model Comparison

```
library(loo)
loo_compare(loo(post_poisson), loo(post_nb)) # warnings about high Pareto k values
```

```
##           elpd_diff se_diff
## model2         0.0        0.0
## model1 -3386.5      510.8
```

```
loo_list <- list(loo(post_poisson, moment_match = TRUE), loo(post_nb, moment_match = TRUE))
loo_compare(loo_list)
```

```
##           elpd_diff se_diff
## model2         0.0        0.0
## model1 -3400.9      515.9
```

```
loo_model_weights(loo_list)
```

```
## Method: stacking
## -----
##           weight
## model1 0.000
## model2 1.000
```

# A Binomial Model for Romney vs Obama in 2012

```
poll <- readRDS("GooglePoll.rds") # WantToWin is coded as 1 for Romney and 0 for Obama
collapsed <- filter(poll, !is.na(WantToWin)) %>%
  group_by(Region, Gender, Urban_Density, Age, Income) %>%
  summarize(Obama = sum(grepl("Obama", WantToWin)), n = n()) %>%
  na.omit
glimpse(collapsed)
```

```
## Rows: 516
## Columns: 7
## Groups: Region, Gender, Urban_Density, Age [143]
## $ Region      <fct> MIDWEST, MIDWEST, MIDWEST, MIDWEST, MIDWEST, MIDWEST, MIDWEST, MID...
## $ Gender      <fct> Female, Female, Female, Female, Female, Female, Female, Female, Fe...
## $ Urban_Density <fct> Rural, Rural, Rural, Rural, Rural, Rural, Rural, Rural, Rural, Rur...
## $ Age         <ord> 18-24, 18-24, 25-34, 25-34, 35-44, 45-54, 45-54, 45-54, 55-64, 55-...
## $ Income      <ord> "25,000-49,999", "50,000-74,999", "25,000-49,999", "50,000-74,999"...
## $ Obama       <int> 6, 5, 4, 3, 8, 10, 8, 0, 0, 18, 4, 10, 1, 6, 21, 13, 1, 1, 0, 40, ...
## $ n          <int> 12, 7, 7, 4, 15, 28, 12, 1, 1, 44, 8, 22, 3, 8, 31, 19, 1, 1, 1, 5...
```

# Prior Predictive Distribution in Symbols

- Here is how McElreath does many hierarchical binomial models
- Suppose a categorical predictor  $x_k$  has  $K$  levels

$$\sigma \sim \mathcal{E}(r)$$

$$\forall k : \beta_k \sim \mathcal{N}(m_k, \sigma)$$

$$\forall k : \mu_k = \frac{1}{1 + e^{-\beta_k}}$$

$$\forall k : y_k \sim \text{Binomial}(n_k, \mu_k)$$

- Aggregating Bernoulli random variables with a common success probability to binomial random variables is much more computationally efficient

# Calling stan

```
stan_data <- with(collapsed, list(N = length(Obama), y = Obama, n = n,  
                                  J = c(nlevels(Region), nlevels(Gender),  
                                        nlevels(Urban_Density), nlevels(Age), nlevels(Income))  
                                  Region = as.integer(Region), Gender = as.integer(Gender),  
                                  Urban_Density = as.integer(Urban_Density),  
                                  Age = as.integer(Age), Income = as.integer(Income),  
                                  prior_only = FALSE, m = rep(0, 5), r = rep(1, 5)))  
post_binomial <- stan("binomial.stan", data = stan_data)
```